

Multi-Agent Corporate Governance & Document Analytics Engine

Project Overview

The **Multi-Agent Corporate Governance & Document Analytics Engine** is an enterprise-grade AI platform designed to automate the analysis of financial, legal, compliance, ESG (Environmental, Social, and Governance), sustainability, audit, and regulatory documents.

Unlike traditional "Chat with PDF" applications, this platform operates as a fully autonomous multi-agent intelligence system where specialized AI agents collaborate, verify each other's work, perform document retrieval, audit findings, and generate executive-level governance reports.

The platform is designed to simulate the workflow used by global consulting and audit firms such as:

- Deloitte
- PwC
- EY
- KPMG
- McKinsey
- BCG

The system enables organizations to upload thousands of pages of reports, filings, and compliance documents and receive verified, source-backed insights through an interactive enterprise dashboard.

Business Problem

Organizations generate massive amounts of documentation:

- Annual Reports
- Financial Statements
- Sustainability Reports
- ESG Disclosures
- SEC Filings
- Risk Assessments
- Legal Contracts
- Compliance Reports
- Internal Audit Documents


Reviewing these documents manually requires significant time, expertise, and resources.


Traditional RAG systems can answer questions but often suffer from:

- Hallucinations
- Missing evidence
- Lack of verification
- Poor auditability
- Limited reasoning capabilities

This platform addresses these issues by introducing a coordinated network of specialized AI agents that perform extraction, retrieval, auditing, validation, and criticism before producing a final report.

System Architecture





Technology Stack

Package Management

UV (Astral)

The project uses UV as the primary package manager for Python.

Benefits:

- Extremely fast dependency resolution
- Faster environment creation
- Reproducible builds
- Better developer experience than pip

Responsibilities:

- Dependency installation
- Virtual environment management
- Lock file generation
- Build reproducibility

Frontend Architecture

Next.js + React + TypeScript

The frontend acts as an enterprise governance dashboard.

Responsibilities:

- User Authentication
 - Document Upload
 - Agent Monitoring
 - Report Visualization
 - Analytics Dashboards
 - Source Verification
 - Organization Management
-

Login & Authentication Module

Enterprise-grade authentication system supporting:

Features

- JWT Authentication
- Refresh Tokens
- RBAC (Role-Based Access Control)
- Multi-Tenant Organizations

User Roles

Compliance Officer

Responsible for:

- Regulatory compliance monitoring
- Governance review
- Policy enforcement

Financial Auditor

Responsible for:

- Financial statement validation
- Revenue verification
- Risk assessment

Legal Analyst

Responsible for:

- Contract review
- Legal clause extraction
- Regulatory interpretation

ESG Reviewer

Responsible for:

- Sustainability reporting
 - Carbon emissions analysis
 - Climate risk evaluation
-

Dashboard Module

The central command center of the platform.

Displays

Uploaded Documents

Annual Report.pdf
SEC Filing.pdf
Sustainability Report.xlsx

Processing Status

Queued
Running
Completed
Failed

Active Agent Jobs

Extraction Agent
Audit Agent
Search Agent
Critic Agent

Analytics Widgets

- Document Statistics
 - Agent Performance
 - Processing Costs
 - Token Usage
 - Success Rates
-

Upload Center

Supports enterprise document ingestion.

Supported Formats

- PDF
- XLSX
- CSV
- DOCX

Workflow

```
graph TD;
  Upload --> Validation;
  Validation --> Storage;
  Storage --> Queue_Creation[Queue Creation];
  Queue_Creation --> Processing_Start[Processing Start];
```

Features:

- Drag and Drop Upload
- Multiple File Upload
- Progress Tracking
- File Validation
- Retry Support

Backend Architecture

FastAPI

FastAPI serves as the AI orchestration backend.

Responsibilities:

- API endpoints
- Authentication
- Agent execution
- Workflow management

- Queue creation
- Database interactions

Benefits

- Asynchronous architecture
 - High throughput
 - Excellent LLM integration
 - Fast performance
-

NestJS Gateway

Acts as an enterprise gateway layer.

Responsibilities:

- API aggregation
- Authentication middleware
- Organization routing
- Request logging
- Rate limiting

Benefits:

- Enterprise scalability
 - Microservice readiness
 - Better separation of concerns
-

Background Processing Layer

Redis

Redis acts as:

- Message Broker
 - Queue Backend
 - Cache Layer
 - Session Store
-

BullMQ

BullMQ manages long-running AI tasks.

Why BullMQ?

Processing large reports may require:

- OCR
- Table extraction
- Embedding generation
- Multi-agent analysis

These tasks may take minutes.

BullMQ ensures:

- Non-blocking processing
 - Retries
 - Job recovery
 - Fault tolerance
-

Job Lifecycle

```
Job Created
↓
Queued
↓
Worker Assigned
↓
Processing
↓
Completed
```

Example:

```
Job #1981

Status:
Queued

Progress:
72%

Estimated Time:
2 Minutes
```

Document Processing Pipeline

Stage 1: File Ingestion Agent

Responsible for understanding uploaded files.

Input Types

- PDFs
 - Excel Files
 - CSV Files
 - Word Documents
-

Extracted Information

Financial Data

Revenue
Profit
Expenses
Assets
Liabilities
Cash Flow

ESG Data

Carbon Emissions
Energy Usage
Water Consumption

Legal Data

Clauses
Obligations
Penalties
Regulations

Compliance Data

```
Violations
Policies
Risk Controls
```

Stage 2: Text Processing

Pipeline:

```
Raw Text
↓
Cleaning
↓
Chunking
↓
Metadata Extraction
↓
Embedding
```

Stage 3: Vector Indexing

All document chunks are stored in:

pgvector

or

Pinecone

Metadata includes:

```
{
  "page": 56,
  "document": "Annual Report",
  "section": "Revenue Analysis"
}
```

Multi-Agent Architecture

The core innovation of the platform.

Router Agent

Purpose

Master decision-making agent.

Receives user requests such as:

Analyze ESG compliance risks.

Determines:

Need Retrieval
Need ESG Analysis
Need Audit
Need Verification

Routes tasks to specialized agents.

Models

- GPT-5
 - Claude
 - Other high-reasoning models
-

Extraction Agent

Purpose

Extract structured information from documents.

Models:

- Qwen
- Llama
- Ollama-hosted models

Produces:

```
{  
  "company": "ABC Corp",  
  "revenue": "$2.1B",  
  "emissions": "12%"  
}
```

Responsibilities:

- Entity Extraction
 - Financial Extraction
 - Table Understanding
 - Metadata Creation
-

Search Agent

Purpose

Advanced Retrieval-Augmented Generation.

Pipeline:

```
Chunk Retrieval  
↓  
Embedding Search  
↓  
Hybrid Search  
↓  
Reranking  
↓  
Context Building
```

Technologies:

- pgvector
 - Pinecone
 - FlashRank
 - Cohere Rerank
-

Audit Agent

Purpose

Cross-verifies extracted information.

Example:

Extraction Agent:

Revenue = \$2.1B

Audit Agent checks:

Page 34
Page 102
Income Statement
Cash Flow Statement

Verifies consistency.

ESG Analysis Agent

Specialized sustainability reviewer.

Analyzes:

- Carbon emissions
- Climate commitments
- Net-zero targets
- Sustainability disclosures
- ESG risks

Produces:

ESG Score
Risk Score
Compliance Rating

Verification Agent

Validates evidence quality.

Responsibilities:

- Evidence matching
- Confidence scoring
- Citation verification
- Source validation

Output:

Claim Supported
Claim Weak
Claim Unsupported

Critic Agent

The most important component.

Purpose:

- Detect hallucinations
- Challenge assumptions
- Verify calculations
- Validate evidence

Example:

Claim:

Emissions reduced 25%

Evidence:

Source indicates 12%

Result:

Claim Rejected

This dramatically increases trustworthiness.

LangGraph Workflow

LangGraph provides stateful orchestration.

Workflow:

```
Router
↓
Extraction
↓
Search
↓
Audit
↓
Verification
↓
Critic
↓
Final Report
```

Benefits:

- Shared State
 - Agent Memory
 - Retry Logic
 - Conditional Routing
 - Human Approval Nodes
-

Hybrid RAG System

A production-grade retrieval system.

Chunking

Documents are split into:

- Semantic Chunks
 - Financial Sections
 - Legal Sections
 - ESG Sections
-

Embeddings

Generated using modern embedding models.

Stored in:

pgvector
or
Pinecone

Hybrid Search

Combines:

Dense Retrieval

Captures semantic meaning.

Sparse Retrieval

Captures exact keyword matches.

Results are merged.

Reranking

Top results pass through:

- FlashRank
- Cohere Rerank

Ensuring the most relevant evidence is selected.

Agent Monitoring Dashboard

Provides complete transparency.

Displays:

```
Agent Status
Execution Time
Cost
Retries
Tokens Used
Confidence Score
```

Example:

```
Router Agent      Completed
Extraction Agent  Completed
Audit Agent       Running
Critic Agent      Waiting
```

Source Verification Interface

Every claim must be traceable.

Example:

```
Claim:
Revenue Increased 18%

Evidence:
Annual Report

Page:
56

Confidence:
98%
```

Users can click evidence and jump directly to the supporting source.

Report Generation System

Produces executive-ready reports.

Sections include:

Executive Summary

High-level findings.

Financial Analysis

Revenue trends.

Profitability.

Cash flow observations.

Compliance Findings

Policy violations.

Regulatory risks.

Audit observations.

ESG Assessment

Carbon emissions.

Sustainability goals.

Climate compliance.

Risk Analysis

Financial risks.

Governance risks.

Operational risks.

Recommendations

Actionable next steps.

Evidence References

Every conclusion linked to supporting sources.

Database Design

PostgreSQL Tables

users

Stores user accounts.

organizations

Stores company workspaces.

documents

Stores uploaded documents.

jobs

Stores processing jobs.

reports

Stores generated reports.

agent_runs

Stores execution history.

audit_logs

Stores verification actions.

Vector Database Schema

document_chunks

Stores chunked content.

embeddings

Stores vector embeddings.

metadata

Stores:

Page Number
Section
Document Type
Timestamp

Human-in-the-Loop Review

Enterprise-grade verification layer.

Allows auditors to:

- Approve findings
- Reject findings
- Edit reports
- Request re-analysis

This creates a collaborative AI workflow rather than a fully autonomous black-box system.

Key Engineering Skills Demonstrated

This project showcases advanced industry-level capabilities:

AI & Agent Systems

- Multi-Agent AI
- LangGraph
- Agent Routing
- Agent Memory
- Critic Agents
- Verification Agents

Retrieval Systems

- Hybrid RAG
- Vector Search
- Reranking
- Embedding Pipelines

Backend Engineering

- FastAPI
- NestJS
- Redis
- BullMQ
- Async Processing

Database Engineering

- PostgreSQL
- pgvector
- Pinecone
- Supabase

Frontend Engineering

- Next.js
- React
- TypeScript
- Data Visualization
- Real-Time Dashboards

Enterprise Features

- RBAC
- Multi-Tenancy
- Audit Trails

- Monitoring
 - Human-in-the-Loop Validation
-

Final Outcome

The Multi-Agent Corporate Governance & Document Analytics Engine is not a simple document chatbot. It is a fully featured enterprise AI platform that combines autonomous agents, advanced retrieval systems, governance analytics, audit verification, ESG intelligence, and executive reporting into a production-grade solution.

The platform demonstrates how modern AI systems can move beyond basic question answering and evolve into trustworthy, auditable, and explainable decision-support systems suitable for financial institutions, consulting firms, regulators, auditors, legal teams, and corporate governance professionals.